

Evaluating Tensorflow Parallel Mechanisms for Training of Deep Learning in a HPC cluster

RAYAN RADDATZ DE MATOS, KENICHI BRUMATI, and MARCELO CARDOSO OLIVEIRA GULART, Federal University of Rio Grande do Sul, Brazil

Put the abstract here.

Additional Key Words and Phrases: Deep Learning, HPC, Parallel Training, Distributed Training

ACM Reference Format:

Rayan Raddatz de Matos, Kenichi Brumati, and Marcelo Cardoso Oliveira Gulart. 2025. Evaluating Tensorflow Parallel Mechanisms for Training of Deep Learning in a HPC cluster. 1, 1 (December 2025), ?? pages. <https://doi.org/0000001.0000001>

1 INTRODUCTION

In the last decade, Deep Learning has established itself as the state-of-the-art for solving complex problems, such as computer vision and natural language processing. Given the progress in algorithmic precision, there has been an exponential growth in the complexity of neural network architectures and the volume of datasets. This scenario imposes a massive computational load, making the training of robust models on a single processing unit an difficult task in terms of time and resources.

To mitigate these challenges, distributed training has become a standard practice, with an emphasis on the Data Parallelism strategy. In this approach, a complete replica of the model is instantiated on each processing unit, and the global dataset is partitioned into smaller, locally-focused batches. Although effective, this technique introduces critical synchronization challenges, as locally calculated gradients must be aggregated and communicated among all devices at every training step, typically through All-Reduce operations.

Given this fact, in this work we decided to analyze data parallelism for AI training using TensorFlow in a HPC cluster to investigate how the factors related to the application contribute to time reduction in a parallel scenario.

Software and Data Availability. We endeavor to make our analysis reproducible for a better science. We made available a companion material hosted in a public GitHub repository at <https://github.com/kenichi220/comp-sys-perf-analysis-final-work/tree/main>. Our companion contains the source code of this article and the software necessary to handle the created datasets. We also include instructions to run the experiment and figures.

2 BACKGROUND AND EXPERIMENTAL CONTEXT

2.1 Background

TensorFlow is a machine learning system that operates at a large scale and in heterogeneous [?]. It uses dataflow graphs to represent the computation, shared state, and the operations that alter this state.

This architecture maps the nodes of a dataflow graph onto many machines within a cluster and, within a machine, onto multiple computational devices (multicore CPUs, GPUs, and TPUs).

Authors' address: Rayan Raddatz de Matos; Kenichi Brumati; Marcelo Cardoso Oliveira Gulart, Federal University of Rio Grande do Sul, Campus do Vale - Sector 4, Porto Alegre, RS, 91501-970, Brazil.

2025. XXXX-XXXX/2025/12-ART \$15.00
<https://doi.org/0000001.0000001>

This offers flexibility to the developer to experiment with new optimizations and training algorithms, supporting a variety of applications focused on deep neural networks. When it comes to distributed training, TensorFlow offers different strategies with the `tf.distribute.Strategy` API. The TensorFlow documentations adverts that *"For synchronous training on many GPUs on multiple workers, use the `tf.distribute.MultiWorkerMirroredStrategy` with the `Keras Model.fit` or a custom training loop."*[?]. Meaning for a distributed training across different nodes in a cluster there is two options: Using the `MultiWorkerMirroredStrategy` with the `model.fit` or creating a Custom Training Loop by yourself.

2.2 Experimental Context

In High-Performance Computing (HPC) environments, a highly recommended practice to access to shared computational resources is to mediate the access by workload management systems. In our context Slurm (Simple Linux Utility for Resource Management) is utilized. Slurm is an open-source, fault-tolerant, and highly scalable job scheduler, responsible for three main functions: Allocating exclusive or non-exclusive access to resources/nodes for users, providing a framework to initiate, execute, and monitor the work on the set of allocated nodes and managing a queue of pending jobs [?].

Interaction with Slurm is typically performed through shell scripts written in the BASH language. These scripts contain preprocessing directives that specify resource requirements, such as the number of nodes, tasks per node, and execution time limit, enabling the automation and reproducibility of computational experiments.

3 METHODS AND MATERIALS

3.1 Software Management and Reproducibility

Reproducibility of computational experiments is a critical challenge in computer science. To mitigate software dependency issues, the use of the GNU Guix package manager was explored. Guix implements the functional package management paradigm where each package is installed in a unique and immutable directory in the store, derived from a cryptographic hash of its dependencies [?].

However, the efficient training of deep neural networks using TensorFlow requires the use of proprietary NVIDIA hardware acceleration libraries, which are not officially distributed in standard Guix channels due to its strict Free Software policy.

Given this restriction, the strategy adopted in this work is the use of Python Virtual Environment for environment management. Although a functional version of Guix exists on GitHub without the proprietary libraries.

For managing the specific Deep Learning libraries, the python package manager `pip` with version 25.3 was utilized. The exact version of the used library was frozen and documented using the `pip freeze` command, generating a requirements file (`reqs.txt`) that enables the reconstruction of the execution environment used in the experiments.

3.2 Hardware and Software Configuration

The experiments were run using TensorFlow 2.15 as newer versions showed errors when trying to train using more than one node with the `.fit` and `MultiWorkerMirrodStrategy` and python 3.11.2. The NVIDIA drivers version was 550.54.15 and CUDA 12.4. We used up to three computational nodes connected with a network of 1G, each node contains an Intel(R) Core(TM) i7-14700KF processor at 3.40 GHz with 28 threads and 20 cores, 96 GB DDR5 RAM and NVIDIA GeForce RTX 4070 with 8GB memory accelerator.

3.3 Design of Experiments (DoE)

We selected two different Deep Learning models to train: **ResNet50** [?] as a big model with ~25.6M parameters and **EfficientNetB0**[?] as a small model with ~5.3M parameters. The two networks were trained using the TinyImageNet-200, a subset of the ImageNet[?] containing 100000 downsized to a 64×64 colored image divided among 200 classes. We created a full factorial for each model with varying batch size and the number of nodes used. The levels for the batch size factor were 32, 64 and 100 and for the nodes we trained with 1, 2 and 3 nodes. Our output variable was the training time (full training time and time per epoch).

4 CHALLENGES

Multi-node communication faced a bottleneck as the node poti4 was connected with a limited bandwidth of 100 Mbps. Regarding software management, GNU Guix does not natively support proprietary libraries due to its strict policy requiring free and open-source software.

5 RESULTS

6 CONCLUSION

ACKNOWLEDGEMENTS

We would like to express our gratitude to Professor Dr. Lucas Mello Schnorr for his guidance and the valuable insights shared regarding performance analysis. We also would like to thank the PCAD at INF/UFRGS for making infrastructure and hardware used in the experiments available.

REFERENCES

- [] Martín Abadi, Paul Barham, Jianmin Chen, Zhifeng Chen, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Geoffrey Irving, Michael Isard, Manjunath Kudlur, Josh Levenberg, Rajat Monga, Sherry Moore, Derek G. Murray, Benoit Steiner, Paul Tucker, Vijay Vasudevan, Pete Warden, Martin Wicke, Yuan Yu, and Xiaoqiang Zheng. 2016. TensorFlow: A System for Large-Scale Machine Learning. In *Proceedings of the 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI '16)*. USENIX Association.
- [] Ludovic Courtès. 2013. Functional Package Management with Guix. In *Proceedings of the 6th European Lisp Symposium (ELS)*. 15–28. <https://arxiv.org/abs/1305.4584>
- [] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*. Ieee, 248–255.
- [] TensorFlow Documentation. 2023. Distributed training with Keras. <https://www.tensorflow.org/tutorials/distribute/keras>. Accessed: 2025-11-30.
- [] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2015. Deep Residual Learning for Image Recognition. *CoRR* abs/1512.03385 (2015). arXiv:1512.03385 <http://arxiv.org/abs/1512.03385>
- [] Mingxing Tan and Quoc V. Le. 2019. EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. *CoRR* abs/1905.11946 (2019). arXiv:1905.11946 <http://arxiv.org/abs/1905.11946>
- [] Andy B. Yoo, Morris A. Jette, and Mark Grondona. 2003. SLURM: Simple Linux Utility for Resource Management. In *Job Scheduling Strategies for Parallel Processing*. Springer, 44–60. https://doi.org/10.1007/10968987_3